

# Reinforcement Learning and Control

Felipe José Bravo Márquez

December 2, 2024

# Introduction

- In supervised learning, algorithms try to make their outputs mimic the labels  $y$  given in the training set <sup>1</sup>.
- The labels give an unambiguous “right answer” for each of the inputs  $x$ .
- In contrast, for many sequential decision making and control problems, it is very difficult to provide this type of explicit supervision to a learning algorithm.
- For example, if we have just built a four-legged robot and are trying to program it to walk.
- Initially we have no idea what the “correct” actions to take are to make it walk
- Hence, we don't know how to provide explicit supervision for a learning algorithm to try to mimic.

---

<sup>1</sup>These slides are based on [Ng, 2000].

# Introduction

- In the reinforcement learning framework, we provide our algorithms only a reward function.
- This function indicates to the learning agent when it is doing well, and when it is doing poorly.
- In the four-legged walking example, the reward function might give the robot positive rewards for moving forwards, and negative rewards for either moving backwards or falling over.
- It will then be the learning algorithm's job to figure out how to choose actions over time so as to obtain large rewards.

# Introduction

- Reinforcement learning has been successful in applications as diverse as:
  - 1 autonomous helicopter flight
  - 2 robot legged locomotion
  - 3 cell-phone network routing
  - 4 marketing strategy selection
  - 5 factory control
  - 6 efficient web-page indexing
- Our study of reinforcement learning will begin with a definition of the Markov decision processes (MDPs).
- MDPs provide the formalism in which RL problems are usually posed.

# Markov Decision Process (MDP)

- A Markov Decision Process is a tuple:

$$(S, A, \{P_{SA}\}, \gamma, R)$$

where

- $S$  is a set of states. (For example, in autonomous helicopter flight,  $S$  might be the set of all possible positions and orientations of the helicopter.)
- $A$  is a set of actions. (For example, the set of all possible directions in which you can push the helicopter's control sticks.)

# Markov Decision Process (MDP)

- $P_{sa}$  are the state transition probabilities. For each state  $s \in S$  and action  $a \in A$ ,  $P_{sa}$  is a distribution over the state space, i.e., it gives the distribution over what states we will transition to if we take action  $a$  in state  $s$ .

$$\sum_{s'} P_{sa}(s') = 1, \quad P_{sa}(s') \geq 0$$

- $\gamma \in [0, 1)$  is a discount factor.
- $R : S \rightarrow \mathcal{R}$  is a reward function.

# Markov Decision Process (MDP)

The dynamics of an MDP proceeds as follows:

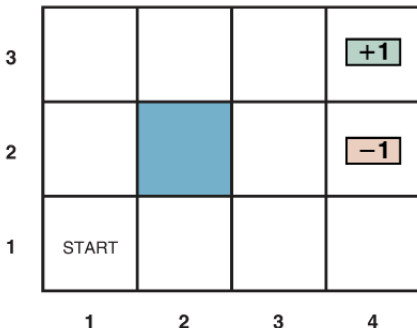
- We start in some state  $s_0$ , and get to choose some action  $a_0 \in A$  to take in the MDP.
- As a result of our choice, the state of the MDP randomly transitions to some successor state  $s_1$ , drawn according to  $s_1 \sim P_{s_0 a_0}$ .
- Then, we get to pick another action  $a_1$ .
- As a result of this action, the state transitions again, now to some  $s_2 \sim P_{s_1 a_1}$ .
- We then pick  $a_2$ , and so on. . . .

Pictorially, we can represent this process as follows:

$$s_0 \xrightarrow{a_0} s_1 \xrightarrow{a_1} s_2 \xrightarrow{a_2} s_3 \xrightarrow{a_3} \dots$$

# Example

- Suppose that an agent is situated in the  $4 \times 3$  environment as shown in the Figure below:

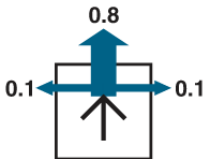


- Beginning in the start state, it must choose an action at each time step.
- There 11 states ( $S$ ) and four actions  $A = \{N, S, E, W\}$ .
- The interaction with the environment terminates when the agent reaches one of the goal states, marked +1 or -1.



# Example

- The “intended” outcome occurs with probability 0.8, but with probability 0.2 the agent moves at right angles to the intended direction:



- For example:

$$P_{(3,1)\mathcal{N}}((3,2)) = 0.8$$

$$P_{(3,1)\mathcal{N}}((4,1)) = 0.1$$

$$P_{(3,1)\mathcal{N}}((2,1)) = 0.1$$

$$P_{(3,1)\mathcal{N}}((3,3)) = 0$$

... etc

# Total Payoff

- Upon visiting the sequence of states  $s_0, s_1, \dots$  with actions  $a_0, a_1, \dots$ , our **total payoff** is given by:

$$R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

$$0 \leq \gamma < 1$$

- Our goal in reinforcement learning is to choose actions over time so as to maximize the expected value of the total payoff:

$$E[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots]$$

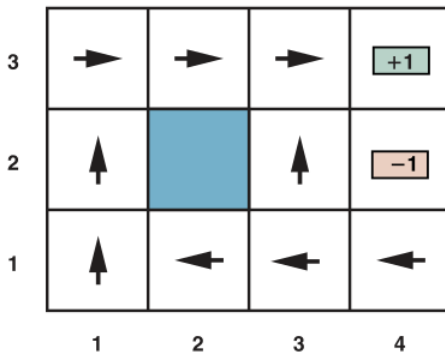
- Note that the reward at timestep  $t$  is discounted by a factor of  $\gamma^t$ .
- Thus, to make this expectation large, we would like to accrue positive rewards as soon as possible (and postpone negative rewards as long as possible).
- In economic applications where  $R(\cdot)$  is the amount of money made,  $\gamma$  also has a natural interpretation in terms of the interest rate (where a dollar today is worth more than a dollar tomorrow).

# Example

- A collision with a wall results in no movement.
- Transitions into the two terminal states have reward  $+1$  and  $-1$ , respectively.  
 $R((4, 3)) = +1$   
 $R((4, 2)) = -1$   
 $R(S) = -0.02$  for all other states.
- All other transitions have a reward of  $-0.02$
- This is common in navigation tasks to charge the robot for fuel consumption and to (to avoid the robot wasting time).
- We will also assume that the robot stops when it reaches states  $(4, 3)$  or  $(4, 2)$ .
- These terminal states can be modelled by adding an extra state in which transitions from these terminal states are made with probability 1 and the agent loops forever in this new state with no more rewards.
-

# Policy

- A policy is any function  $\pi : S \rightarrow A$ , mapping states to actions.
- We say that we are executing some policy  $\pi$  if, whenever we are in state  $s$ , we take action  $a = \pi(s)$ .
- The “optimal policy” for the previous example would be:



- So when you execute this policy this will maximize your expected value of the total payoffs.
- It is worth noting that the action of going west (left) from state  $(3, 1)$  is not trivial, since going north would bring the agent closer to the goal.
- However, entering state  $(3, 2)$  adds the danger of ending up in state  $(4, 2)$ .
- Therefore, the optimal solution favored the longer and less risky path.

# Value Function

- For any policy  $\pi$ , we define the value function  $V^\pi : \mathcal{S} \rightarrow \mathbb{R}$  as:

$$V^\pi(s) = \mathbb{E} \left[ R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \cdots \mid s_0 = s, \pi \right].$$

- $V^\pi(s)$  is the expected sum of discounted rewards (or total payoff) starting in state  $s$  and executing policy  $\pi$ .

# Value Function

- The following examples shows a pretty bad policy  $\pi$  than seems to be heading to -1 rather than +1 in most cases:

|   |   |   |   |    |
|---|---|---|---|----|
| 3 | → | → | → | +1 |
| 2 | ↓ |   | → | -1 |
| 1 | → | → | ↑ | ↑  |

- The value function  $V^\pi(s)$  for this policy is as follows:

|   |      |      |      |       |
|---|------|------|------|-------|
| 3 | .52  | .73  | .77  | +1    |
| 2 | -.90 |      | -.82 | -1    |
| 1 | -.88 | -.87 | -.85 | -1.00 |

# References I



Ng, A. (2000).

Cs229 lecture notes.

*CS229 Lecture notes*, 1(1):1–3.